

# Movielens\_project

April 5, 2021

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib import style
%matplotlib inline
```

```
[8]: #Import the three datasets
movies = pd.read_csv("movies.dat", sep="::", names=['MovieID', 'Title',
↳ 'Genres'])
ratings = pd.read_csv("ratings.dat", sep="::", names=['UserID', 'MovieID',
↳ 'Rating', 'Timestamp'])
users = pd.read_csv("users.dat", sep="::", names=['UserID', 'Gender', 'Age',
↳ 'Occupation', 'Zip-code'])
```

<ipython-input-8-3f2fb1081e27>:2: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

```
movies = pd.read_csv("movies.dat", sep="::", names=['MovieID', 'Title',
'Genres'])
```

<ipython-input-8-3f2fb1081e27>:3: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

```
ratings = pd.read_csv("ratings.dat", sep="::", names=['UserID', 'MovieID',
'Rating', 'Timestamp'])
```

<ipython-input-8-3f2fb1081e27>:4: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

```
users = pd.read_csv("users.dat", sep="::", names=['UserID', 'Gender', 'Age',
'Occupation', 'Zip-code'])
```

```
[9]: movies.head()
```

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy

2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

```
[10]: ratings.head()
```

```
[10]:   UserID  MovieID  Rating  Timestamp
0      1      1193      5  978300760
1      1       661      3  978302109
2      1       914      3  978301968
3      1      3408      4  978300275
4      1      2355      5  978824291
```

```
[11]: users.head()
```

```
[11]:   UserID  Gender  Age  Occupation  Zip-code
0      1      F    1      10      48067
1      2      M   56      16      70072
2      3      M   25      15      55117
3      4      M   45       7      02460
4      5      M   25      20      55455
```

```
[12]: #Create a new dataset [Master_Data] with the following columns MovieID Title
      ↳UserID Age Gender Occupation Rating.
      #(i) Merge two tables at a time. (ii) Merge the tables using two primary keys
      ↳MovieID & UserId
      userRatings = pd.merge(users, ratings, on=['UserID'])
      movieRatings = pd.merge(movies, ratings, on=['MovieID'])
```

```
[13]: userRatings.head()
```

```
[13]:   UserID  Gender  Age  Occupation  Zip-code  MovieID  Rating  Timestamp
0      1      F    1      10      48067      1193      5  978300760
1      1      F    1      10      48067       661      3  978302109
2      1      F    1      10      48067       914      3  978301968
3      1      F    1      10      48067      3408      4  978300275
4      1      F    1      10      48067      2355      5  978824291
```

```
[14]: movieRatings.head()
```

```
[14]:   MovieID      Title      Genres  UserID  Rating  \
0      1  Toy Story (1995)  Animation|Children's|Comedy      1      5
1      1  Toy Story (1995)  Animation|Children's|Comedy      6      4
2      1  Toy Story (1995)  Animation|Children's|Comedy      8      4
3      1  Toy Story (1995)  Animation|Children's|Comedy      9      5
4      1  Toy Story (1995)  Animation|Children's|Comedy     10      5

      Timestamp
```

```

0 978824268
1 978237008
2 978233496
3 978225952
4 978226474

```

```
[15]: Master_Data = pd.merge(userRatings, movieRatings, on=['UserID', 'MovieID', 'Rating'])
```

```
[16]: Master_Data.head()
```

```
[16]:
```

	UserID	Gender	Age	Occupation	Zip-code	MovieID	Rating	Timestamp_x \
0	1	F	1	10	48067	1193	5	978300760
1	1	F	1	10	48067	661	3	978302109
2	1	F	1	10	48067	914	3	978301968
3	1	F	1	10	48067	3408	4	978300275
4	1	F	1	10	48067	2355	5	978824291

	Title	Genres \
0	One Flew Over the Cuckoo's Nest (1975)	Drama
1	James and the Giant Peach (1996)	Animation Children's Musical
2	My Fair Lady (1964)	Musical Romance
3	Erin Brockovich (2000)	Drama
4	Bug's Life, A (1998)	Animation Children's Comedy

	Timestamp_y
0	978300760
1	978302109
2	978301968
3	978300275
4	978824291

```
[17]: Master_Data = Master_Data[['MovieID', 'Title', 'UserID', 'Age', 'Gender', 'Occupation', 'Rating']]
```

```
[18]: Master_Data.head()
```

```
[18]:
```

	MovieID	Title	UserID	Age	Gender	\
0	1193	One Flew Over the Cuckoo's Nest (1975)	1	1	F	
1	661	James and the Giant Peach (1996)	1	1	F	
2	914	My Fair Lady (1964)	1	1	F	
3	3408	Erin Brockovich (2000)	1	1	F	
4	2355	Bug's Life, A (1998)	1	1	F	


	Occupation	Rating
0	10	5
1	10	3

2	10	3
3	10	4
4	10	5

```
[19]: #Explore the datasets using visual representations (graphs or tables), also
      ↪ include your comments on the following:
      #User Age Distribution
      #User rating of the movie "Toy Story"
      #Top 25 movies by viewership rating
      #Find the ratings for all the movies reviewed by for a particular user of user
      ↪ id = 2696
```


```
[20]: #User Age Distribution
users['Age'].hist(bins=50)
plt.xlabel('Age')
plt.ylabel('Population')
plt.show
```

```
[20]: <function matplotlib.pyplot.show(close=None, block=None)>
```




```
[23]: #User rating of the movie "Toy Story"
toy = Master_Data[Master_Data.Title == "Toy Story (1995)"]
plt.plot(toy.groupby("Age")["MovieID"].count())
toy.groupby("Age")["MovieID"].count()
```

```
[23]: Age
      1      112
      18     448
      25     790
      35     423
      45     143
      50     108
      56      53
      Name: MovieID, dtype: int64
```



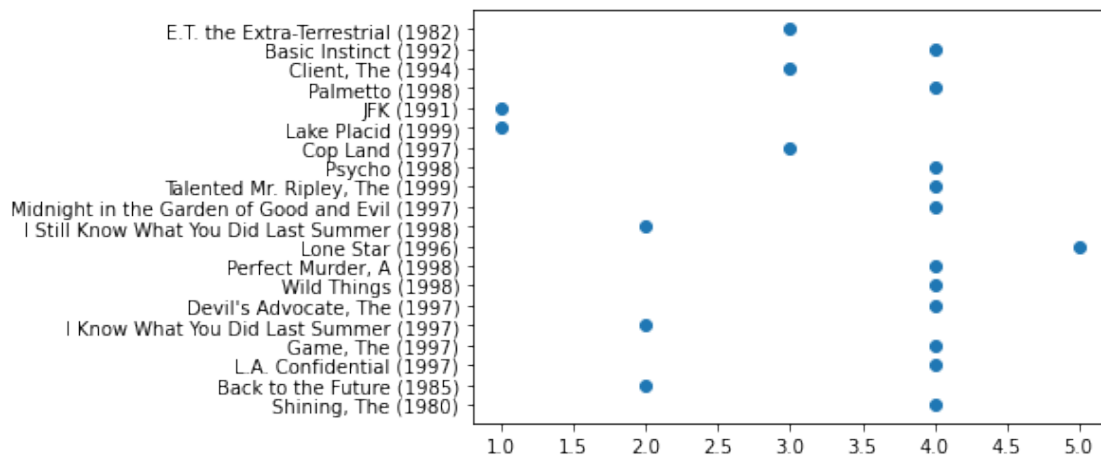
```
[35]: #Top 25 movies by viewership rating
top = Master_Data.groupby("Title").size().sort_values(ascending=False).head(25)
plt.ylabel("Title")
plt.xlabel("Viewership Count")
top.plot(kind='barh')
```

```
[35]: <AxesSubplot:xlabel='Viewership Count', ylabel='Title'>
```



```
[37]: #Find the ratings for all the movies reviewed by for a particular user of user_
      ↪ id = 2696
res = Master_Data[Master_Data.UserID == 2696]
plt.scatter(y=res.Title, x=res.Rating)
```

```
[37]: <matplotlib.collections.PathCollection at 0x26f6b109b50>
```



```
[ ]: #Feature Engineering:
      #Use column genres:
      #Find out all the unique genres (Hint: split the data in column genre making a
      ↪ list)
      #process the data to find out only the unique categories of genres)
      #Create a separate column for each genre category with a one-hot encoding ( 1
      ↪ and 0) whether or not the movie belongs to that genre.
      #Determine the features affecting the ratings of any particular movie.
```

```
#Develop an appropriate model to predict the movie ratings
```

```
[38]: #Use column genres:  
#Find out all the unique genres (Hint: split the data in column genre making a  
↳list)  
val = movies.Genres.str.split("|")
```

```
[39]: val.head()
```

```
[39]: 0      [Animation, Children's, Comedy]  
1      [Adventure, Children's, Fantasy]  
2              [Comedy, Romance]  
3              [Comedy, Drama]  
4              [Comedy]  
Name: Genres, dtype: object
```

```
[41]: res_col = []  
for v in val:  
    for i in v:  
        if i not in res_col:  
            res_col.append(i)
```

```
[42]: res_col.append("Gender")  
res_col.append("Age")  
res_col.append("Rating")  
df = pd.DataFrame(columns=res_col)  
df.head()
```

```
[42]: Empty DataFrame  
Columns: [Animation, Children's, Comedy, Adventure, Fantasy, Romance, Drama,  
Action, Crime, Thriller, Horror, Sci-Fi, Documentary, War, Musical, Mystery,  
Film-Noir, Western, Gender, Age, Rating]  
Index: []  
  
[0 rows x 21 columns]
```

```
[48]: res = Master_Data.merge(movies, on = ['MovieID'],  
↳how="left")[["Genres", "Rating", "Gender", "Age"]]
```

```
[51]: #Create a separate column for each genre category with a one-hot encoding ( 1  
↳and 0) whether or not the movie belongs to that genre.  
for index, row in res.head(20000).iterrows():  
    tmp = row.Genres.split("|")  
    for i in tmp:  
        df.loc[index, i] = 1  
        df.loc[index, "Gender"] = res.loc[index, "Gender"]  
        df.loc[index, "Age"] = res.loc[index, "Age"]
```

```

df.loc[index,"Rating"] = res.loc[index,"Rating"]
df.loc[index,df.columns[~df.columns.
→isin(tmp+["Gender","Rating","Age"])]]= 0

```

```
[52]: df.head()
```

```

[52]:
      Animation Children's Comedy Adventure Fantasy Romance Drama Action \
19999      0      0      1      0      0      0      1      1
0      0      0      0      0      0      0      1      0
1      1      1      0      0      0      0      0      0
2      0      0      0      0      0      1      0      0
3      0      0      0      0      0      0      1      0

      Crime Thriller ... Sci-Fi Documentary War Musical Mystery Film-Noir \
19999      0      0 ...      0      0      0      0      0
0      0      0 ...      0      0      0      0      0
1      0      0 ...      0      0      1      0      0
2      0      0 ...      0      0      1      0      0
3      0      0 ...      0      0      0      0      0

      Western Gender Age Rating
19999      0      M  25      3
0      0      F   1      5
1      0      F   1      3
2      0      F   1      3
3      0      F   1      4

```

[5 rows x 21 columns]

```

[55]: #Develop an appropriate model to predict the movie ratings
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

x = df[df.columns[~df.columns.isin(["Rating"])]
y = df.Rating

x_train, x_test, y_train, y_test = train_test_split(x, y, random_state = 0)

number = LabelEncoder()
x_train.Gender = number.fit_transform(x_train["Gender"].astype("str"))
x_test.Gender = number.fit_transform(x_test["Gender"].astype("str"))
y_train = number.fit_transform(y_train.astype("int"))
y_test = number.fit_transform(y_test.astype("int"))

```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\generic.py:5168:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.



Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
self[name] = value
```

```
[57]: #SVM

from sklearn.svm import SVC
svm_model_linear = SVC(kernel = 'linear', C = 1).fit(x_train, y_train)
svm_predictions = svm_model_linear.predict(x_test)

# model accuracy for x_test
accuracy = svm_model_linear.score(x_test, y_test)

# creating a confusion matrix
cm = confusion_matrix(y_test, svm_predictions)
accuracy
```

```
[57]: 0.3412
```

```
[64]: cm
```

```
[64]: array([[134, 20, 62, 58, 13],
        [162, 15, 143, 135, 28],
        [322, 62, 428, 436, 87],
        [418, 60, 552, 541, 149],
        [233, 61, 317, 405, 159]], dtype=int64)
```

```
[59]: #KNN

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 7).fit(x_train, y_train)

# accuracy on X_test
accuracy = knn.score(x_test, y_test)

# creating a confusion matrix
knn_predictions = knn.predict(x_test)
cm = confusion_matrix(y_test, knn_predictions)

accuracy
```

```
[59]: 0.3052
```

```
[65]: cm
```

```
[65]: array([[134, 20, 62, 58, 13],
            [162, 15, 143, 135, 28],
            [322, 62, 428, 436, 87],
            [418, 60, 552, 541, 149],
            [233, 61, 317, 405, 159]], dtype=int64)
```

```
[61]: #Naive Bayes classifier

from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB().fit(x_train, y_train)
gnb_predictions = gnb.predict(x_test)

# accuracy on X_test
accuracy = gnb.score(x_test, y_test)

# creating a confusion matrix
cm = confusion_matrix(y_test, gnb_predictions)

accuracy
```

```
[61]: 0.2554
```

```
[63]: from sklearn.linear_model import LinearRegression
lr = LinearRegression().fit(x_train, y_train)
lr_predictions = lr.predict(x_test)

# accuracy on X_test
accuracy = lr.score(x_test, y_test)

accuracy
```

```
[63]: 0.02635611144303962
```

```
[ ]:
```